

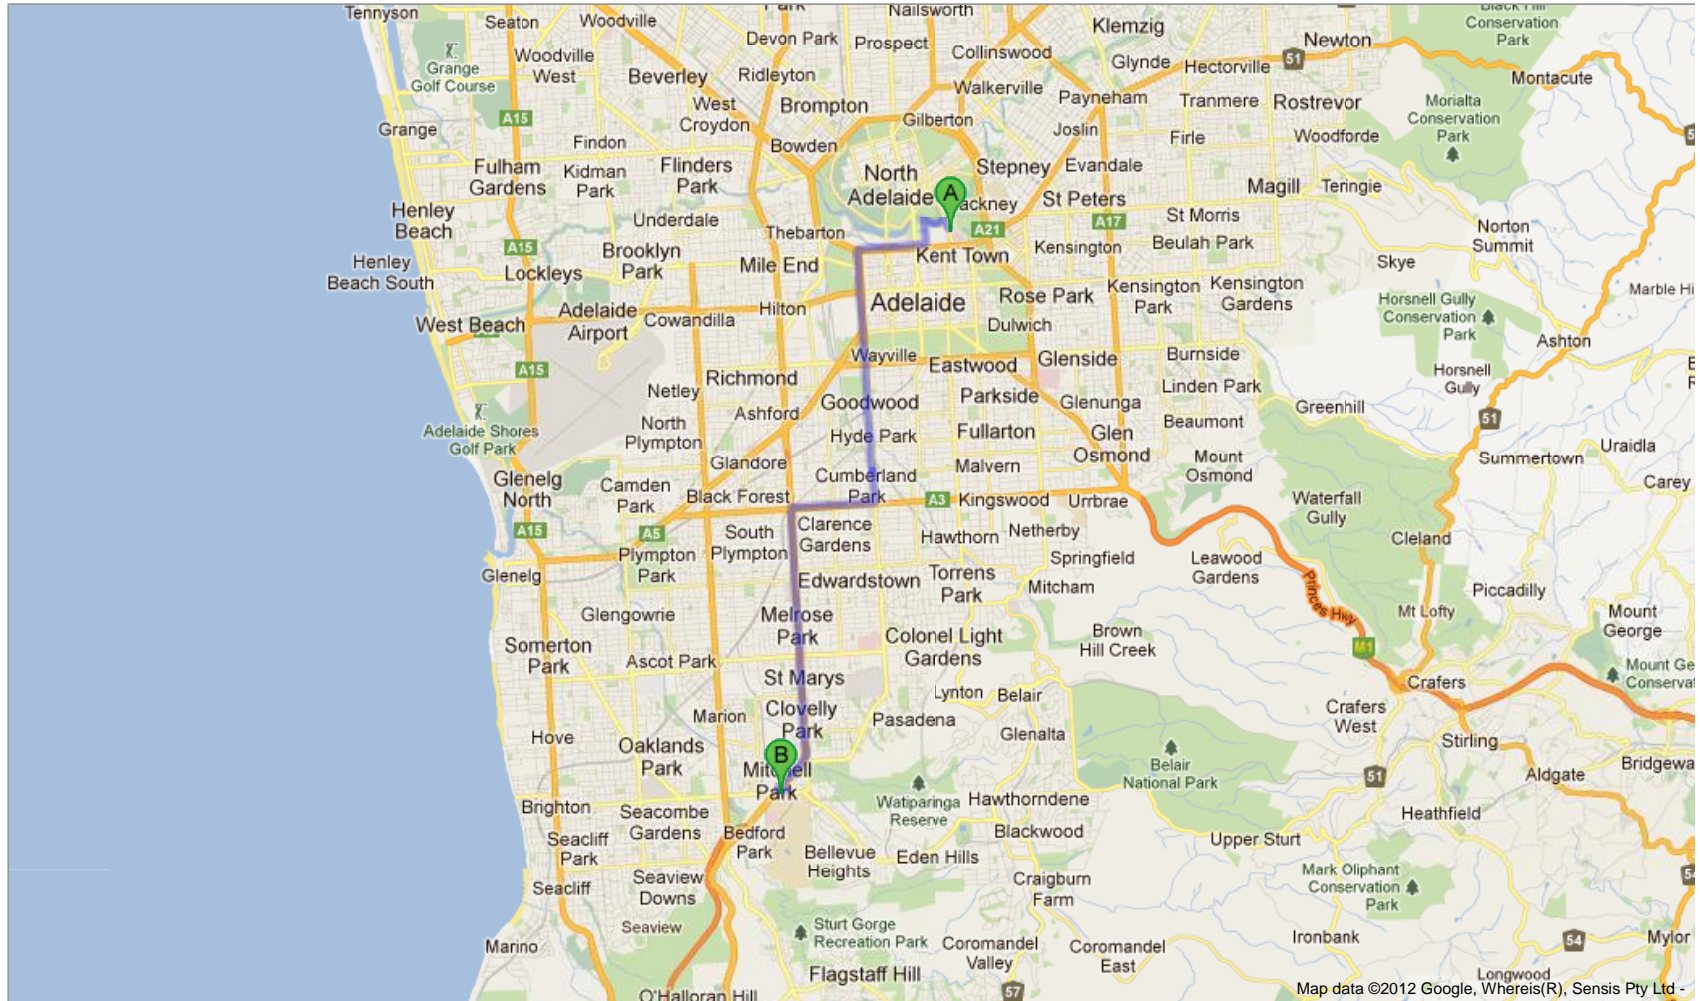
Algorithm and Data Structure Analysis (ADSA)

Shortest Paths

Problem

- Computation of shortest path is one of the classical problems.
- Classical application is route planning.

Uni Adelaide – Flinders Uni



Problem Statement

Given a directed graph $G=(V,E)$ and a cost function $c : E \rightarrow R$ on the edges.

Given a path $p = (e_1, e_2, \dots, e_k)$ consisting of k edges the cost of the path is

$$c(p) = \sum_{i=1}^k c(e_i)$$

A **shortest path** from a node s to a node v is a path of minimal cost among all possible paths from s to v .

Problem Statement

Single-source shortest path problem:

Compute for a given node s of V a shortest path to any other node in V (if it exists).

We assume that edge weights are non-negative.

Why non-negative edge costs?

If a path from s to v contains a negative cycles then a shortest path does not exist (is not defined).




Fig 10.1 Mehlhorn/Sanders

Simple shortest path for non-negative edge costs

If edge costs are non-negative and v is reachable from s then a shortest path P from s to v exists. P can be chosen to be simple (cycle-free).

Properties of subpaths

Lemma: Subpaths of a shortest path are also shortest paths.

Proof (by contradiction):

- Assume that the path P is a shortest path from s to v .
- Assume that a subpath from a to b is not a shortest path from a to b



- This implies that there is a shorter path from a to b
- We can use this path to obtain a shorter path from s to v .
- Contradiction to P is shortest path from s to v .




Dijkstra's Algorithm

- Remember BFS for computing all shortest paths in an unweighted graph.
- In iteration i , we computed all shortest paths having i edges.
- Dijkstra's algorithm obtains in iteration i a shortest path to the node of the i th smallest distance from s .
- We can represent all shortest paths from a node s by a tree rooted at s .


Dijkstras Algorithm

We call a node u unscanned if no shortest path from s to u has been found so far




Dijkstras Algorithm

Make unscanned node u scanned that would get the minimal tentative distance among all unscanned nodes.




Dijkstras Algorithm

Make unscanned node u scanned that would get the minimal tentative distance among all unscanned nodes.



Updating

We may update a previous path from s to v if we find a shorter path




Procedure $relax(e = (u, v) : Edge)$

if $d[u] + c(e) < d[v]$ **then** $d[v] := d[u] + c(e); \quad parent[v] := u$

Dijkstras Algorithm

Consider all edges leaving u and update distances using relax.



Dijkstras Algorithm

Dijkstra's Algorithm


declare all nodes unscanned and initialize d and $parent$

while there is an unscanned node with tentative distance $< +\infty$ **do**

$u :=$ the unscanned node with minimal tentative distance

 relax all edges (u, v) out of u and declare u scanned

Example




Nodes for which a shortest path has been computed




Nodes that have already been reached


Example




Example




Example




Example




Example




Example




Example




Example




Example




Example



Example



Example



Correctness

Theorem 10.5: Dijkstra's algorithm solves the single-source shortest path problem for graphs with nonnegative edge costs.

Proof:

We show two steps:

- All nodes reachable from s are scanned after termination.
- When a node v becomes scanned then the shortest path from s to v is obtained.

Correctness

Claim: All nodes reachable from s are scanned after termination.

Proof (by contradiction):

- Consider a shortest path $p=(s=v_1, v_2, \dots, v_k=v)$ from s to v
- Let $i>1$ be minimal such **that v_i is unscanned**
- Implies node v_{i-1} has been scanned.
- When v_{i-1} is scanned $d[v_i]$ is set to $d[v_{i-1}]+c(v_{i-1},v_i) < \infty$.
- Hence, **v_i must be scanned** as only nodes u with $d[u]=\infty$ stay unscanned. **Contradiction to v_i is unscanned.**



Correctness

Claim: When a node v becomes scanned then the shortest path from s to v is obtained.

Proof (by contradiction):

- Denote by $\mu[v]$ the length of a shortest path from s to v .
- Consider the first point in time t when v has been scanned and $d[v] > \mu[v]$ holds. (we assumed the shortest path from s to v is not obtained when v becomes scanned)
- Consider a shortest path $p=(s=v_1, v_2, \dots, v_k=v)$ from s to v .
- Let $i>1$ be minimal such that v_i has not been scanned before time t .

Correctness

Proof (continued):

- Node v_{i-1} was scanned before time t which implies $\mu[v_{i-1}] = d[v_{i-1}]$.
- When v_i is scanned $d[v_i] = d[v_{i-1}] + c(v_{i-1}, v_i) = \mu[v_{i-1}] + c(v_{i-1}, v_i) = \mu[v_i]$.
- We have $d[v_i] = \mu[v_i] \leq \mu[v_k] < d[v_k]$
- When $i=k$, we have $d[v_k] < d[v_k]$, a contradiction.



Implementation

- Store all unscanned reached nodes in an addressable priority queue Q (using tentative distances as key values)

Pseudocode Dijkstra

Function *Dijkstra*($s : \text{NodeId}$) : $\text{NodeArray} \times \text{NodeArray}$

$d = \langle \infty, \dots, \infty \rangle : \text{NodeArray}$ **of** $\mathbb{R} \cup \{\infty\}$

$\text{parent} = \langle \perp, \dots, \perp \rangle : \text{NodeArray}$ **of** NodeId

$\text{parent}[s] := s$

$Q : \text{NodePQ}$

$d[s] := 0$; $Q.\text{insert}(s)$

while $Q \neq \emptyset$ **do**

$u := Q.\text{deleteMin}$

foreach $\text{edge } e = (u, v) \in E$ **do**

if $d[u] + c(e) < d[v]$ **then**

$d[v] := d[u] + c(e)$

$\text{parent}[v] := u$

if $v \in Q$ **then** $Q.\text{decreaseKey}(v)$

else $Q.\text{insert}(v)$

return (d, parent)



// returns (d, parent)
// tentative distance from root

// self-loop signals root
// unscanned reached nodes

// we have $d[u] = \mu(u)$

// relax

// update tree



Runtime

- Initialization (arrays, priority queue) takes time $O(n)$.
- Every reachable node is inserted and removed once from Q.
- At most **n deleteMin and insert operations**.
- Each node is scanned at most once and each edge is relaxed at most once.
- Implies at most **m decreaseKey** operations.

Total runtime

$$T_{\text{Dijkstra}} = O\left(m \cdot T_{\text{decreaseKey}}(n) + n \cdot (T_{\text{deleteMin}}(n) + T_{\text{insert}}(n))\right)$$

Runtime

Runtime depends on implementation of priority queue.

Original (Dijkstra 1959):

- Maintain the number of reached unscanned nodes.
- An array d storing the distances and an array storing for each node whether it is reached or unscanned.
- Insert and decreaseKey take time $O(1)$
- DeleteMin takes time $O(n)$
- Total Runtime: $O(m+n^2)$

Improvements:

- Binary Heaps: $O((m+n) \log n)$
- Fibonacci Heaps: $O(m + n \log n)$